

Retrieval-Augmented Generation (RAG) System for Sanskrit Text Question Answering

1. Problem Statement

Classical Sanskrit texts are rich in philosophical and religious knowledge but are difficult to query using traditional keyword-based search systems.

The objective of this project is to build a **Retrieval-Augmented Generation (RAG)** pipeline that enables **semantic question answering** over Sanskrit documents using modern NLP techniques.

The system must:

- Accept questions in **Sanskrit or English**
 - Retrieve relevant Sanskrit passages
 - Generate **grounded answers only from retrieved context**
 - Avoid hallucinations
 - Run on **CPU-only hardware**
-

2. Overall System Design

The system follows a standard **RAG architecture**:



This separation ensures **interpretability, modularity, and correctness**.

3. Data Processing & Chunking

3.1 Source Data

- Sanskrit verses stored in text format
- File: processed_chunks.txt

3.2 Chunking Strategy

- Text is split into **logical verse-based chunks**
- Each chunk preserves semantic coherence
- Chunk size chosen to balance:
 - Semantic completeness
 - Embedding efficiency

3.3 Storage

- Chunks stored as:
 - NumPy array: chunk_texts.npy
 - Used later for retrieval mapping

4. Embedding Model (Core Technical Choice)

Model Used

```
bash  
  
sentence-transformers/all-MiniLM-L6-v2
```

Why this model?

- Lightweight (384-dimensional embeddings)
- CPU-efficient
- Good semantic performance
- Supports multilingual semantic similarity

Embedding Details

- Input: Sanskrit / English text
- Output: 384-d float32 vector
- Normalization: handled internally by SentenceTransformers

5. Vector Indexing with FAISS

Index Type

```
nginx
```

```
IndexFlatL2
```

Why FAISS?

- Efficient similarity search
- Scales well for medium-sized corpora
- CPU-compatible

Index Construction

- All chunk embeddings added to FAISS index
- Stored persistently as:

```
python
```

```
faiss_index.bin
```

Distance Metric

- L2 (Euclidean distance)
-

6. Retriever Module

Function

```
python
```

```
retrieve_chunks(query, top_k=3)
```

Steps

1. Encode query using embedding model
2. Perform FAISS search
3. Retrieve top-k nearest chunk indices
4. Map indices back to Sanskrit text chunks

Key Parameters

Parameter	Value
top_k	3
embedding_dim	384
distance_metric	L2

Design Decision

- Using **semantic similarity**, not keyword match
- Allows English queries to retrieve Sanskrit content

7. Context Construction

Purpose

To limit LLM input size and ensure relevance.

Implementation

```
python
```

```
MAX_CONTEXT_CHARS = 1200
```

- Chunks concatenated until limit reached
- Prevents context overflow
- Ensures answer grounding

8. Generator Module (Answering Logic)

Input

- Retrieved Sanskrit context
- Original user query

Behavior

- Answers only if information is present in context
- Returns fallback message otherwise
- Does not hallucinate
- Preserves original language (English / Sanskrit)

Key Design Choice

No generative paraphrasing

Answers are extractive / context-aware

Why this matters

- Ensures factual correctness
 - Ideal for classical texts
 - Avoids misleading outputs
-

9. Query Language Handling

Query Language Output Language

English English

Sanskrit Sanskrit

This behavior emerges naturally from:

- Multilingual embeddings
 - Sanskrit-only context
 - Non-translational generation logic
-

10. Execution Environment

Hardware

- CPU only
- No GPU required

Dependencies

```
bash

faiss-cpu
sentence-transformers
numpy
torch
```

Reasoning

- Ensures portability

- Suitable for academic evaluation
 - No cloud dependency
-

11. Evaluation & Results

Observations

- Correct retrieval of semantically relevant verses
- Accurate grounding of answers
- Graceful failure when answer is absent
- Stable performance across multiple queries

Example

Query (Sanskrit):

भक्तः कः?

Answer:

भक्तः सः यः परमेश्वरं निरन्तरं स्मरति।

12. Limitations

- No abstractive summarization
 - Limited to provided corpus
 - No fine-tuned Sanskrit LLM
 - No translation layer (by design)
-

13. Future Improvements

- Sanskrit-specific embedding models
- Larger corpus indexing

- Hybrid keyword + semantic retrieval
 - Abstractive answer generation
 - Translation + bilingual output mode
-

14. Conclusion

This project demonstrates a **robust, interpretable, and efficient RAG pipeline** for Sanskrit texts.

It successfully bridges ancient language processing with modern NLP techniques while maintaining correctness and computational efficiency.